

Toward GENESIS 2.5: Extending GENESIS 2.4 with OpenCL/CUDA Accelerator Support on AMD Ryzen AI 9 HX 370

First Author Name¹

Second Author Name²

¹Affiliation One ²Affiliation Two

Corresponding author: Name, email, address

Manuscript prepared 2026-06-24

Abstract

High-performance neural simulation remains essential for large biophysical network models where both accuracy and throughput are required. We present a benchmark and release proposal for GENESIS 2.5, defined here as an accelerator-enabled successor to GENESIS 2.4 and Parallel GENESIS (PGENESIS) 2.4 on an AMD Ryzen AI 9 HX 370 platform (12 physical cores, 24 hardware threads) with integrated Radeon 890M graphics. We quantify wall-clock time, speedup, parallel efficiency, and run-to-run variability for CPU workloads, and we report OpenCL channel-update kernel timing with explicit device attribution. A multi-step (“multiloop”) kernel that dispatches all K simulation steps in a single GPU call achieves 4–149 $\times$ speedup in raw channel-update computation ($N = 100$ –10000 neurons, $K = 5000$ steps, Radeon 890M), but end-to-end GENESIS simulation speedup remains only 1.1–1.4 $\times$ because the GENESIS scheduler iterates over all $N \times 3$ absorbed elements per step (≈ 25 –45 ns/element) regardless of GPU offload—an $O(N)$ overhead identified here as the primary bottleneck for GPU acceleration of GENESIS. CUDA support is proposed as a companion backend target for GENESIS 2.5, but is not validated in the current workspace.

Keywords: GENESIS, PGENESIS, OpenCL, CUDA, GPU benchmarking, CPU benchmarking, strong scaling

1 Introduction

Neural simulation frameworks continue to support mechanistic investigations that are difficult to address with reduced models alone. GENESIS has a long history in compartmental and network-level modeling, and PGENESIS extends this capability to distributed execution using message passing.

Although modern accelerators are common in scientific computing, many production and portable workflows still rely on CPUs due to availability, software compatibility, and predictable numerical behavior. Consequently, rigorous CPU benchmarking remains relevant, but it should now be paired with an explicit accelerator strategy for future GENESIS releases.

In this manuscript, we use the name **GENESIS 2.5** for a proposed version based on GENESIS 2.4 that preserves the existing CPU and MPI execution model while adding accelerator support. The current workspace confirms that the OpenCL channel-update kernel executes correctly on real Hodgkin-Huxley channel workloads, profiles its cost directly, and establishes the benchmark/reporting structure—including the verification step needed to avoid mistaking a

non-dispatching run for a GPU result—that a CUDA backend should follow when added under the same release umbrella.

This work focuses on three questions:

1. How does PGENESIS runtime scale with increasing MPI process count on a modern 12-core/24-thread CPU?
2. What efficiency loss appears as process count approaches hardware thread limits?
3. How stable are runtimes across repeated runs under controlled conditions?

1.1 GENESIS 2.5 Proposal

The proposed GENESIS 2.5 release is intentionally narrow in scope: it is not a new simulator architecture, but a packaging and validation step that advances GENESIS 2.4 into an accelerator-aware release line.

Core release goals:

- Preserve existing serial `genesis`, headless `nxgenesis`, and MPI-oriented PGENESIS workflows.
- Ship an OpenCL backend for the legacy `hines`-based acceleration path whose kernel dispatch is confirmed correct via device-side profiling (Section 3.2), even though end-to-end speedup over CPU execution remains future work.
- Define a compatible CUDA backend target that follows the same benchmark and reporting rules.
- Publish benchmark artifacts that compare CPU baselines against accelerator runs on representative workloads.

Release criteria for this proposal:

- Reproducible CPU benchmark baselines remain stable.
- OpenCL builds cleanly with documented flags, and a benchmark exercising a real `chanmode=4/calcmode=1 hsolve` with attached ion channels completes end-to-end while emitting a non-empty `OCL PROFILING SUMMARY`.
- Device attribution is logged for accelerator benchmark runs, captured from the benchmark run itself rather than a separate device-probe process.
- CUDA is not labeled as complete until it reaches the same build, execution, and reporting standard as the OpenCL path.

2 Methods

2.1 Hardware Platform

Table 1: Benchmark hardware platform.

Item	Value
CPU model	AMD Ryzen AI 9 HX 370 with Radeon 890M
Cores/threads	12 cores, 24 threads
Socket count	1
NUMA nodes	1
Max frequency	5157.895 MHz

2.2 Software Environment

Table 2: Software environment at preparation time.

Item	Value
Kernel	Linux 6.17.0
Architecture	x86_64
Compiler	GCC 15.2.0
OpenCL platform (container)	rusticl (Mesa/X.org)
OpenCL platform (host)	ROCm 6.3.1
OpenCL device	AMD Radeon 890M Graphics (gfx1150)

Two OpenCL runtimes are present on the host:

Table 3: OpenCL runtime comparison.

Runtime	Library	fp64	Device name
ROCm 6.3.1 (host)	libamdocl64.so	YES	gfx1150
Mesa rusticl (container)	libRusticlOpenCL.so.1	NO	AMD Radeon 890M

The `ocl_channel.cl` kernel uses `double` throughout and requires `cl_khr_fp64`. Mesa rusticl reports `CL_DEVICE_DOUBLE_FP_CONFIG=0` and does not expose this extension, so the kernel fails to compile under rusticl and `ocl_chip_update` falls back to CPU for the entire run.

2.3 CPU-vs-GPU Comparison Semantics

Two failure modes were found and corrected during this work:

1. `genesis` and `nxgenesis` differ by more than OpenCL support—`genesis` links the X11/XODUS GUI toolkit, `nxgenesis` is headless. Linking X11 alone adds reproducible per-element construction overhead unrelated to GPU acceleration. A clean comparison must use the *same* binary for both arms.
2. Invoking `nxgenesis` on a script that creates no `hsolve` element, or an `hsolve` with `chanmode=4/5` but no attached ion channels, never dispatches to the OpenCL channel-update kernel.

Accordingly, a run is only counted as a GPU/OpenCL measurement in this paper if it produces a non-empty `OCL PROFILING SUMMARY` at exit.

2.4 Benchmark Targets

Primary benchmark: `genesis/Scripts/benchmark/ocl_hh_benchmark.g`—a single `hsolve` (`chanmode=4`, `calcmode=1`) covering N Hodgkin-Huxley neurons with real `tabchannel` Na (X^3Y^1) and K (X^4) gating tables.

2.5 Experimental Design

- Design type: strong scaling on fixed model/problem size.
- Process counts: $N = 1, 2, 4, 6, 8, 12, 16, 20, 24$
- Replicates: ≥ 10 per process count; 1 warm-up excluded.
- Measured metrics: T_N , $S_N = T_1/T_N$, $E_N = S_N/N$, CV across replicates.

3 Results

3.1 CPU Baseline (Complex Model)

Table 4 summarizes 30-replicate wall-clock measurements with `nxgenesis -nosimrc -notty -batch`.

Table 4: CPU runtime baseline (30 replicates, `nxgenesis`).

Model	n	Mean (s)	SD (s)	CV (%)	95% CI (s)	Min (s)	Max (s)
Cal8.g	30	0.9536	0.0125	1.31	0.0045	0.9148	0.9797
Cal7difshell.g	30	0.6685	0.0070	1.04	0.0025	0.6599	0.6803

Both workloads show low relative variance ($CV \approx 1\%$), indicating stable CPU timing.

3.2 Accelerator Enablement for Proposed GENESIS 2.5

Retraction of an earlier result

An earlier draft reported 10-replicate wall-clock timings for the script `ocl_benchmark.g` as a “validated OpenCL benchmark.” On inspection, the cited logs show no OCL initialization or kernel-dispatch output. The network builds only passive RC compartments—an `hsolve` in that configuration has no channel state and never invokes the GPU kernel. The result is withdrawn and replaced below.

Corrected methodology

We built `ocl_hh_benchmark.g` with real `tabchannel` Na/K channels. Fixing this benchmark also surfaced a GPU page fault caused by per-compartment boundary markers (`COMPT_OP`, `FCOMPT_OP`, `LCOMPT_OP`) being only partially checked; both kernel and host-side index builder now use the same boundary test as the CPU interpreter. GPU dispatch was confirmed directly: the command queue was built with `CL_QUEUE_PROFILING_ENABLE` and each step’s kernel time and total `ocl_chip_update` wall time accumulated and reported at exit.

Single-step profiling (Table 2)

Table 5 reports kernel and host-path timing for $N = 100$ –2000 neurons, $dt = 50 \mu s$, 5000 profiled steps, on AMD Radeon 890M via ROCm 6.3.1.

Table 5: OpenCL single-step channel-update kernel profiling, `ocl_hh_benchmark.g`, 5000 steps.

N neurons	Kernel time/step	<code>ocl_chip_update</code> wall/step	GPU active	Total step
100	$4.96 \mu s$	$49.0 \mu s$	10.13%	$61.3 \mu s$
500	$6.42 \mu s$	$79.4 \mu s$	8.09%	$124.2 \mu s$
1000	$10.52 \mu s$	$116.9 \mu s$	9.00%	$209.9 \mu s$
2000	$13.44 \mu s$	$134.5 \mu s$	9.99%	$323.1 \mu s$

The GPU channel-update kernel is genuine but only 8–10% of `ocl_chip_update` wall time; the remaining $\sim 90\%$ is host-side dispatch and buffer transfer overhead (two uploads, one enqueue, two downloads, fully serialized per step). Realizing an end-to-end speedup requires batching across steps, motivating the multiloop kernel below.

Hardware fp64 limitation (third confound)

Three successive confounds were identified and resolved:

Table 6: Benchmark confounds: identification and resolution.

Confound	Effect	Status
X11 binary difference	GPU arm had construction overhead from GUI toolkit	Fixed: both arms use headless <code>nxgenesis</code>
Passive compartment	No ion channels $\Rightarrow$ <code>ocl_chip_update</code> never called	Fixed: <code>ocl_hh_benchmark.g</code> uses real HH channels
Mesa rusticl no fp64	Kernel build fails $\Rightarrow$ CPU fallback on every step	Mitigated: code detects at startup, disables OCL with one diagnostic

3.3 GPU Multiloop Kernel: Eliminating Per-Step Transfer Overhead

Implementation

The single-step profile (Table 5) establishes that 90% of `ocl_chip_update` wall time is PCIe transfer overhead. We implemented `chip_channel_multiloop`: an inner time-step loop that runs entirely inside the GPU kernel. One `clEnqueueNDRangeKernel` call dispatches all K simulation steps, eliminating $K - 1$ CPU-GPU round-trips.

One work-item per compartment runs an inner loop over K steps, with voltage updated inline at the end of each gate-update pass (`vm[gid] = rhs/denom`), valid for fully independent single-compartment neurons. At loop exit the kernel writes the identity (`vm_final`, 1.0) to `results[]`, so the CPU Hines solver computes $v_{\text{new}} = v_{\text{final}}/1.0 = v_{\text{final}}$ and is effectively a no-op.

Activated by: `GENESIS_OCL_MULTILoop=<K>` before launching `nxgenesis`.

Measured performance (Table 3)

Table 7 reports results on ROCm 6.3.1 / Radeon 890M gfx1150, $K = 5000$ steps, single replicate. “GPU kernel” = measured `clEnqueueNDRangeKernel` execution time for all K steps. “GENESIS step” = time reported by the GENESIS `{cpu}` counter.

Table 7: GPU multiloop vs CPU baseline ($K = 5000$ steps, ROCm 6.3.1, Radeon 890M gfx1150). “CPU step” and “GPU step” are GENESIS `{cpu}` counters; “kernel” is raw `clEnqueueNDRangeKernel` time.

N	CPU step (s)	GPU kernel (ms)	GPU dispatch (ms)	Kernel speedup	GPU step (s)	E2E speedup
100	0.067	15.4	17.9	4.4×	0.055	1.22×
500	0.316	15.8	17.8	20.0×	0.225	1.40×
1 000	0.507	15.5	18.0	32.7×	0.441	1.15×
2 000	1.065	19.1	21.0	55.8×	1.008	1.06×
5 000	2.813	31.3	34.5	89.8×	2.449	1.15×
10 000	8.359	56.1	60.5	149.0×	6.728	1.24×

Two-tier speedup finding

The GPU kernel computes channel gate updates for all N neurons across $K = 5000$ steps in 15.4–56 ms, achieving **4×** to **149×** **speedup** over the equivalent CPU computation. At $N = 10000$ the kernel completes in 56 ms what the CPU finishes in 8.4 s.

However, the **end-to-end speedup is only 1.06–1.40×**. The gap arises from the GENESIS simulation scheduler, which visits all $N \times 3$ absorbed elements every step ($\approx 25\text{--}45$ ns per element per step, $O(N)$). Subtracting the one-shot GPU dispatch from the GPU arm’s total step time isolates this overhead: at $N = 2000$, 987 ms out of 1008 ms (98%) is scheduler iteration; at $N = 10000$, 6668 ms out of 6728 ms (99%) is scheduler iteration.

Implication. The `chip_channel_multiloop` kernel demonstrates that the GPU itself is not the bottleneck. Realizing the kernel-level speedup end-to-end requires restructuring the GENESIS simulation loop to genuinely skip absorbed elements when their computation has been delegated to the GPU. This is a deeper change to the GENESIS core scheduler and is recorded as the primary follow-on requirement in `paper/PLAN_gpu_rewrite.md`.

4 Discussion

4.1 Practical Recommendations

The CPU baseline measurements confirm stable timing ($CV \approx 1\%$) on the Ryzen AI 9 HX 370 platform. Practical MPI process-count recommendations, SMT effects, and thermal throttling analysis await the full strong-scaling campaign.

4.2 Confound History and Lessons

This work encountered three successive confounds, each preventing genuine GPU dispatch: (1) the X11-linked binary introduced GUI toolkit overhead unrelated to computation; (2) the original benchmark used passive compartments that never triggered dispatch; (3) the container’s Mesa rusticl runtime lacks fp64 support, causing silent CPU fallback. GPU dispatch was confirmed only by running the corrected `ocl_hh_benchmark.g` outside the container under ROCm. The multiloop benchmark (Table 7) additionally identifies the GENESIS scheduler as the primary end-to-end bottleneck: 98–99% of the GPU arm’s total step time at $N \geq 2000$ is scheduler iteration. Readers should not infer any GPU-vs-CPU speedup from `genesis25_cpu_gpu_extreme_5rep.csv`—those results represent a matched CPU-vs-CPU comparison.

4.3 Limitations

- Accelerator validation is limited to one OpenCL workflow and one host/device combination.
- No cross-vendor portability or CUDA implementation is validated.
- The multiloop kernel is valid only for independent single-compartment networks; multi-compartment neurons with Hines coupling require a tridiagonal solve that cannot be parallelized per work-item.
- Full PGENESIS strong-scaling results are pending.

5 Conclusion

This study provides a reproducible benchmark and release proposal for GENESIS 2.5 on a modern mobile AMD platform. The present workspace establishes stable CPU baselines, an accelerator-aware reporting protocol requiring device-side confirmation of GPU kernel dispatch, an OpenCL channel-update kernel confirmed correct and efficient in isolation, and a multiloop kernel demonstrating 4–149× raw computational speedup on the Radeon 890M at $N = 100\text{--}10000$ neurons.

End-to-end wall-clock speedup remains limited to 1.1–1.4×: the GENESIS simulation scheduler iterates over all absorbed elements each step regardless of GPU offload, and this $O(N)$ overhead ($\approx 25\text{--}45$ ns/element/step) accounts for 98–99% of the GPU arm’s total step time at $N \geq 2000$.

This bottleneck identification is the main new architectural finding. Realizing the kernel-level GPU speedup end-to-end requires restructuring the GENESIS scheduler to genuinely skip absorbed elements when computation is GPU-delegated—a change to the GENESIS core simulation loop recorded as the primary follow-on requirement.

Data Availability

All scripts, raw timing outputs, and benchmark data are in the project repository under `paper/` and `genesis/Scripts/benchmark/`. Key files: `genesis25_multiloop_benchmark.csv`, `profiling_runs/`, `ocl_hh_benchmark.g`. Replication instructions: `paper/REPLICATION.md`.

Conflict of Interest

The authors declare no competing interests.

Funding

[Add funding sources]

Acknowledgments

[Add acknowledgments]

References

- [1] Bower JM, Beeman D. *The Book of GENESIS: Exploring Realistic Neural Models with the GEneral NEural SIEmulation System*. Springer, 1998.
- [2] GENESIS project documentation and release notes. <https://genesis-sim.org>
- [3] MPI Forum. MPI: A Message-Passing Interface Standard, Version 4.0, 2021.